

# Modeling

This notebook combines the three cleaned datasets into one central location and aggregates them before conducting data engineering and running a wide array of models to determine the final top performer and understand the relationships between the features.

## Data Imports

In [1]:

```
# Basics
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import sys
import pickle

# Importing databases using SQL
from sqlalchemy import create_engine

# Model preprocessing and processing
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.compose import make_column_selector, make_column_transformer
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import GridSearchCV
from imblearn.pipeline import Pipeline
from sklearn.base import clone

# Models
from sklearn.dummy import DummyClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.naive_bayes import GaussianNB
from xgboost import XGBClassifier

# Performance evaluation
from sklearn.metrics import f1_score, precision_score, accuracy_score, recall_score
from sklearn.metrics import confusion_matrix, plot_confusion_matrix
from sklearn.inspection import permutation_importance

# Data visualization
import shap

# Options
#pd.options.display.max_rows = 200
pd.options.display.max_columns = 200
%matplotlib inline

# Convenience for working with external src code files
%load_ext autoreload
%autoreload 2
sys.path.insert(1, '../src')
```

```

# Custom functions
from create_target import *
from remove_missing_data import *
from evaluate_model_performance import *
from custom_plots import *
from identify_collinearity import *

# Global constants
RANDOM_STATE = 2021

```

### Import "Protests" dataset

```

In [2]: engine = create_engine('sqlite:///../data/processed/protests.db')
        with engine.begin() as connection:
            df_protests = pd.read_sql('SELECT * FROM protests', con=connection)

# Type casting
df_protests.startdate = pd.to_datetime(df_protests.startdate)

```

### Import "Governments" dataset

```

In [3]: engine = create_engine('sqlite:///../data/processed/governments.db')
        with engine.begin() as connection:
            df_govts = pd.read_sql('SELECT * FROM governments', con=connection)

# Set index to be used on Join Later
df_govts.index = df_govts.year_scode

# Remove unused features
df_govts.drop('year_scode', axis=1, inplace=True)

```

### Join "Protests" and "Governments" datasets

```

In [4]: # Join both dataframes
        df = df_protests.join(df_govts, how='left', on='year_scode')

# Remove entries that don't have corresponding 'government' data
df.dropna(inplace=True)
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 15064 entries, 0 to 15207
Data columns (total 76 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   country                              15064 non-null  object
 1   scode                                15064 non-null  object
 2   region                              15064 non-null  object
 3   protestnumber                       15064 non-null  int64
 4   protesterviolence                   15064 non-null  int64
 5   startdate                           15064 non-null  datetime64[ns]
 6   duration_days                       15064 non-null  int64
 7   participants                        15064 non-null  int64
 8   participants_category               15064 non-null  object
 9   demand_labor-wage-dispute          15064 non-null  int64
10  demand_land-farm-issue             15064 non-null  int64
11  demand_police-brutality             15064 non-null  int64

```

12	demand_political-behavior/process	15064	non-null	int64
13	demand_price-increases/tax-policy	15064	non-null	int64
14	demand_removal-of-politician	15064	non-null	int64
15	demand_social-restrictions	15064	non-null	int64
16	year_scode	15064	non-null	object
17	participants_log	15064	non-null	float64
18	duration_days_log	15064	non-null	float64
19	protestnumber_log	15064	non-null	float64
20	system	15064	non-null	object
21	yrsoffc	15064	non-null	float64
22	finittrm	15064	non-null	float64
23	yrcurnt	15064	non-null	float64
24	termlimit	15064	non-null	float64
25	reelect	15064	non-null	float64
26	multpl	15064	non-null	float64
27	military	15064	non-null	float64
28	defmin	15064	non-null	float64
29	prtyin	15064	non-null	float64
30	execrlc	15064	non-null	object
31	execnat	15064	non-null	float64
32	execrel	15064	non-null	object
33	execage	15064	non-null	float64
34	allhouse	15064	non-null	float64
35	totalseats	15064	non-null	float64
36	oppmajh	15064	non-null	float64
37	oppmajs	15064	non-null	float64
38	legelec	15064	non-null	float64
39	exelec	15064	non-null	float64
40	liec	15064	non-null	float64
41	eiec	15064	non-null	float64
42	mdmh	15064	non-null	float64
43	mdms	15064	non-null	float64
44	ssh	15064	non-null	float64
45	plurality	15064	non-null	float64
46	pr	15064	non-null	float64
47	housesys	15064	non-null	object
48	sensys	15064	non-null	object
49	thresh	15064	non-null	float64
50	cl	15064	non-null	float64
51	gq	15064	non-null	float64
52	gqi	15064	non-null	float64
53	fraud	15064	non-null	object
54	auton	15064	non-null	float64
55	muni	15064	non-null	float64
56	state	15064	non-null	float64
57	author	15064	non-null	float64
58	numvote	15064	non-null	float64
59	oppvote	15064	non-null	float64
60	maj	15064	non-null	float64
61	partyage	15064	non-null	float64
62	herfgov	15064	non-null	float64
63	herfopp	15064	non-null	float64
64	frac	15064	non-null	float64
65	oppfrac	15064	non-null	float64
66	govfrac	15064	non-null	float64
67	tensys_strict	15064	non-null	float64
68	checks	15064	non-null	float64
69	stabs_strict	15064	non-null	float64
70	tenlong_strict	15064	non-null	float64
71	tenshort_strict	15064	non-null	float64
72	polariz	15064	non-null	float64
73	country_govt	15064	non-null	object
74	scode_govt	15064	non-null	object
75	percent	15064	non-null	float64

```
dtypes: datetime64[ns](1), float64(51), int64(11), object(13)
memory usage: 8.8+ MB
```

### Import "Regime Changes" dataset

```
In [5]: # IMPORT REGIME CHANGE DATASET
engine = create_engine('sqlite:///../data/processed/regime_changes.db')
with engine.begin() as connection:
    df_regimes = pd.read_sql('SELECT * FROM regime_changes', con=connection)

# Type conversions
df_regimes.startdate = pd.to_datetime(df_regimes.startdate)
df_regimes.enddate = pd.to_datetime(df_regimes.enddate)
df_regimes.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1296 entries, 0 to 1295
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   country          1296 non-null   object
1   scode             1296 non-null   object
2   startdate        1296 non-null   datetime64[ns]
3   enddate          1296 non-null   datetime64[ns]
4   duration_yrs     1296 non-null   float64
5   xconst           1296 non-null   int64
6   present          1296 non-null   int64
dtypes: datetime64[ns](2), float64(1), int64(2), object(2)
memory usage: 71.0+ KB
```

### QC that country names and country IDs match

```
In [6]: cols = ['scode', 'scode_govt', 'country', 'country_govt']
missing_countries = df.loc[(df.country != df.country_govt)][cols]
missing_countries = missing_countries.drop_duplicates()
display(missing_countries.sort_values(by='scode'))
```

scode	scode_govt	country	country_govt
-------	------------	---------	--------------

### Remove countries that do not contain government data

```
In [7]: scodes_to_remove = missing_countries.scode.unique()
scodes_to_remove_ind = [x in scodes_to_remove for x in df.scode]
df.drop(df.loc[scodes_to_remove_ind].index, axis=0, inplace=True)
```

### Identify countries that are missing from "Regime Changes" dataset

```
In [8]: # ALL countries in union of Protests and Governments
all_countries = df.scode.unique()

# ALL countries in Regimes
regime_countries = df_regimes.scode.unique()

# Loop over all_countries
missing = []
for country in all_countries:
    # Make note of any countries not in Regimes
```



```

27  execnat          11840 non-null float64
28  execrel          11840 non-null object
29  totalseats       11840 non-null float64
30  oppmajh          11840 non-null float64
31  legelec          11840 non-null float64
32  exelec           11840 non-null float64
33  liec             11840 non-null float64
34  eiec             11840 non-null float64
35  gq               11840 non-null float64
36  gqi              11840 non-null float64
37  auton            11840 non-null float64
38  numvote          11840 non-null float64
39  oppvote          11840 non-null float64
40  maj              11840 non-null float64
41  herfgov          11840 non-null float64
42  govfrac          11840 non-null float64
43  tensys_strict    11840 non-null float64
44  checks           11840 non-null float64
45  stabs_strict     11840 non-null float64
46  tenlong_strict   11840 non-null float64
47  tenshort_strict  11840 non-null float64
48  country_govt     11840 non-null object
49  scode_govt       11840 non-null object
50  xconst            11840 non-null float64
51  present           11840 non-null float64
52  next_regime_chg_date 11840 non-null datetime64[ns]
53  days_until_next_regime_chg 11840 non-null float64
dtypes: datetime64[ns](2), float64(43), object(9)
memory usage: 5.0+ MB

```

```

In [11]: # Convert startdate to a float instead of datetime since datetime
# cannot be handled by models but fractional years can
df['startdate'] = df.startdate.dt.year + \
                 df.startdate.dt.month/12 + \
                 df.startdate.dt.day/365

```

```

In [12]: # Convert to Categorical datatypes
df['region'] = df.region.astype('category')
df['system'] = df.system.astype('category')
df['country'] = df.country.astype('category')

```

## Adjust encoding of *xconst*

The *xconst* ("Executive Constraints") ranges from 1 ("Unlimited Authority") to 7 ("Executive Parity"). However, there are also three different placeholder values to represent a period of instability. An "interruption period" is coded as -66. An "interregnum period" is coded as -77. A "transition period" is coded as -88. Clearly, these outliers need to be addressed. They cannot be removed because they are inherently valuable when studying regime changes. There is also no conventional way to "normalize" these values. Lastly, given the strong impact this feature has on the model, it has been determined that a one-hot encoding scheme is not the most valuable. Instead, I created my own encoding after meticulous review of the data dictionary (see page 19).

Instead of these periods being represented by -66, -77, and -88, they are instead represented as -1, -2, and 0, respectively. Keeping in mind how this metric will be interpreted as continuous by the model, I have chosen to rank these time periods on a scale of "less stable" to "more stable", where the least stable is the farthest away from achieving "Executive Parity" (7). In summary:

1. Transition periods, previously encoded as -88, are now encoded as 0.
2. Interruption periods, previously encoded as -66, are now encoded as -1.
3. Interregnum periods, previously encoded as -77, are now encoded as -2.

```
In [13]: print('Before:\n', df.xconst.value_counts())
df.xconst.replace(-66.0, -1, inplace=True) # 'Interruption periods'
df.xconst.replace(-77.0, -2, inplace=True) # 'Interregnum periods'
df.xconst.replace(-88.0, 0, inplace=True) # 'Transition periods'
print('\nAfter:\n', df.xconst.value_counts())
```

Before:

```
7.0      4796
6.0      1994
3.0      1592
5.0      1511
2.0       674
4.0       510
1.0       426
-88.0     161
-77.0     114
-66.0      62
Name: xconst, dtype: int64
```

After:

```
7.0      4796
6.0      1994
3.0      1592
5.0      1511
2.0       674
4.0       510
1.0       426
0.0       161
-2.0      114
-1.0       62
Name: xconst, dtype: int64
```

## Define target

This allows the user to define the target in terms of the amount of time before which a regime transition will occur. For this analysis, it uses 1 year, but other values have also been explored with similar results. The shorter the time period, the lower the model performance - as would be expected.

```
In [14]: DAYS_UNTIL_CHG = 365

target = pd.DataFrame(df['days_until_next_regime_chg'] < DAYS_UNTIL_CHG)
target = target.astype('int')
target.columns = ['target']
```

## Drop unused columns

```
In [15]: drop_cols = ['year_scode', 'scode_govt', 'country_govt', 'startdate',
                    'days_until_next_regime_chg', 'scode', 'participants_category',
                    'participants', 'next_regime_chg_date', 'index', 'duration_days',
                    'present', 'protestnumber']
```


```
model_inputs = df.drop(drop_cols, axis=1)
model_inputs.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 11840 entries, 0 to 15060
Data columns (total 41 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   country                                   11840 non-null  category
1   region                                   11840 non-null  category
2   protesterviolence                        11840 non-null  float64
3   demand_labor-wage-dispute               11840 non-null  float64
4   demand_land-farm-issue                  11840 non-null  float64
5   demand_police-brutality                  11840 non-null  float64
6   demand_political-behavior/process        11840 non-null  float64
7   demand_price-increases/tax-policy        11840 non-null  float64
8   demand_removal-of-politician             11840 non-null  float64
9   demand_social-restrictions               11840 non-null  float64
10  participants_log                         11840 non-null  float64
11  duration_days_log                        11840 non-null  float64
12  protestnumber_log                        11840 non-null  float64
13  system                                   11840 non-null  category
14  yrsoffc                                  11840 non-null  float64
15  finittrm                                 11840 non-null  float64
16  termlimit                                11840 non-null  float64
17  military                                  11840 non-null  float64
18  defmin                                   11840 non-null  float64
19  execnat                                  11840 non-null  float64
20  execrel                                  11840 non-null  object
21  totalseats                               11840 non-null  float64
22  oppmajh                                  11840 non-null  float64
23  legelec                                  11840 non-null  float64
24  exelec                                   11840 non-null  float64
25  liec                                    11840 non-null  float64
26  eiec                                    11840 non-null  float64
27  gq                                       11840 non-null  float64
28  gqi                                       11840 non-null  float64
29  auton                                    11840 non-null  float64
30  numvote                                  11840 non-null  float64
31  oppvote                                  11840 non-null  float64
32  maj                                       11840 non-null  float64
33  herfgov                                  11840 non-null  float64
34  govfrac                                  11840 non-null  float64
35  tensys_strict                            11840 non-null  float64
36  checks                                   11840 non-null  float64
37  stabs_strict                             11840 non-null  float64
38  tenlong_strict                           11840 non-null  float64
39  tenshort_strict                          11840 non-null  float64
40  xconst                                   11840 non-null  float64
dtypes: category(3), float64(37), object(1)
memory usage: 3.6+ MB
```

## Identify and resolve multi-collinearity

```
In [16]: calculate_collinearity(model_inputs, min_threshold=0.5)
```





Features with correlation higher than 0.5:

cc	
pairs	
(protesterviolence, protesterviolence)	1.000000
(yrsoffc, tenlong_strict)	0.780163
(tenshort_strict, yrsoffc)	0.765382
(liec, eiec)	0.674035
(tenlong_strict, tenshort_strict)	0.633976
(liec, totalseats)	0.573247
(liec, maj)	0.553506
(oppvote, numvote)	0.504882

Check for high Variance Inflation Factors (VIFs), indicating problematic collinearity

Note the correlation between *liec* and *eiec* in the above heat map is the highest on the plot. Investigate this relationship alongside other features using VIF analysis.

Note that the VIF threshold of 10 is higher than usual. This is because high multi-collinearity is less of an issue for tree-based models, as this notebook deems the most appropriate. A threshold of 10 ensures that extreme collinearity is addressed while also ensuring that features are not unnecessarily dropped, losing valuable predictive information.

```
In [17]: # Calculate Variance Inflation Factor for input DataFrame
def calc_vif(df_input):
    # Source: Flatiron School course material
    # https://github.com/learn-co-curriculum/dsc-modeling-your-data


    vif = [variance_inflation_factor(df_input.values, i) for i in range(df_input.shape[
    return list(zip(df_input.columns, vif))
```

```
In [18]: # Ignore categoricals for VIF
vif_droppers = ['country', 'region', 'system', 'execrel'] # Categoricals
collinearity_df = model_inputs.drop(vif_droppers, axis=1)

display(calc_vif(collinearity_df))
calculate_collinearity(collinearity_df, min_threshold=0.2, plot=True)
```

```
[('protesterviolence', 1.44868299774804),
 ('demand_labor-wage-dispute', 1.6204214619562236),
 ('demand_land-farm-issue', 1.208810302966097),
 ('demand_police-brutality', 1.3395906398870971),
 ('demand_political-behavior/process', 5.513852371701452),
 ('demand_price-increases/tax-policy', 1.3168762166422847),
 ('demand_removal-of-politician', 1.425511446325062),
 ('demand_social-restrictions', 1.1676559446483674),
 ('participants_log', 9.263351416533459),
 ('duration_days_log', 1.1461953736415431),
 ('protestnumber_log', 3.410186100779725),
 ('yrsoffc', 8.37284616777815),
 ('finittrm', 54.41943320661126),
 ('termlimit', 4.873827265679371),
 ('military', 1.5299962092037278),
 ('defmin', 1.474407369303405),
 ('execnat', 1.3426018999607543),
 ('totalseats', 3.771894688181355),
 ('oppmajh', 1.0624530491055708),
 ('legelec', 1.5828221998186602),
 ('exelec', 1.4835143009886813),
 ('liec', 99.23294309707873),
 ('eiec', 64.8298165175228),
 ('gq', 3.084535900311382),
 ('gqi', 1.8876143244248036),
 ('auton', 1.861315081362755),
 ('numvote', 5.697092005756747),
 ('oppvote', 5.929034892174823),
 ('maj', 23.87346082448925),
 ('herfgov', 1.0692269050561742),
 ('govfrac', 2.3605737922563983),
 ('tensys_strict', 5.422521995236603),
 ('checks', 9.594313096599041),
 ('stabs_strict', 1.4338279978327284),
 ('tenlong_strict', 6.652744071314912),
```

(`'tenshort_strict'`, 5.777808363310377),  
(`'xconst'`, 10.436065434582105)]



Features with correlation higher than 0.2:

cc	
pairs	
( <b>protesterviolence</b> , <b>protesterviolence</b> )	1.000000
( <b>yrsoffc</b> , <b>tenlong_strict</b> )	0.780163
( <b>tenshort_strict</b> , <b>yrsoffc</b> )	0.765382
( <b>liec</b> , <b>eiec</b> )	0.674035
( <b>tenlong_strict</b> , <b>tenshort_strict</b> )	0.633976
...	...
( <b>yrsoffc</b> , <b>execnat</b> )	0.209019
( <b>eiec</b> , <b>tenlong_strict</b> )	0.206541
( <b>oppvote</b> , <b>yrsoffc</b> )	0.204862
( <b>military</b> , <b>oppvote</b> )	0.203807

## pairs

(military, gq) 0.202514

78 rows × 1 columns

Remove *liec*

Reminder of definitions:

- *liec*: legislative index of electoral competitiveness
- *eiec*: executive index of electoral competitiveness

Given the similar nature of these features, collinearity is not surprising. Given the downstream analysis of past models, it was determined that *eiec* has a stronger impact on model performance than *liec*. Drop the latter.


```
In [19]: model_inputs.drop(['liec', 'eiec', 'tenlong_strict', 'tenshort_strict',
                          'finitttrm', 'govfrac', 'oppvote'], axis=1, inplace=True)
```

```
In [20]: # Ignore categoricals for VIF
vif_droppers = ['country', 'region', 'system', 'execrel'] # Categoricals
collinearity_df = model_inputs.drop(vif_droppers, axis=1)

display(calc_vif(collinearity_df))
calculate_collinearity(collinearity_df, min_threshold=0.2, plot=True)
```

```
[('protesterviolence', 1.4284165923417442),
 ('demand_labor-wage-dispute', 1.5775585041959777),
 ('demand_land-farm-issue', 1.198318309118901),
 ('demand_police-brutality', 1.3113263252488623),
 ('demand_political-behavior/process', 5.182846671644404),
 ('demand_price-increases/tax-policy', 1.29445052765119),
 ('demand_removal-of-politician', 1.3739765526630505),
 ('demand_social-restrictions', 1.1555870540176596),
 ('participants_log', 8.347742330546884),
 ('duration_days_log', 1.1440221138508397),
 ('protestnumber_log', 3.3552785319353164),
 ('yrsoffc', 2.80425614997517),
 ('termimit', 3.7274704558416283),
 ('military', 1.498776804304648),
 ('defmin', 1.433254004090043),
 ('execnat', 1.2781958948352632),
 ('totalseats', 2.588238035302926),
 ('oppmajh', 1.0557336248410882),
 ('legelec', 1.5578114828080951),
 ('exelec', 1.4671253709952061),
 ('gq', 2.905367989418863),
 ('gqi', 1.8010361525378558),
 ('auton', 1.7968805145686535),
 ('numvote', 3.225797973005871),
 ('maj', 11.970029851031912),
 ('herfgov', 1.0596117867189228),
 ('tensys_strict', 4.957215632157678),
 ('checks', 5.870319376019616),
```

```
('stabs_strict', 1.3645299784376301),
('xconst', 9.271719398976682)]
```



Features with correlation higher than 0.2:

cc	
pairs	
(protesterviolence, protesterviolence)	1.000000
(maj, checks)	0.436117
(gq, tensys_strict)	0.432867
(checks, tensys_strict)	0.420018
(tensys_strict, xconst)	0.416724
(auton, totalseats)	0.402273
(yrsoffc, military)	0.398234
(exelec, legelec)	0.364723
(demand_political-behavior/process, demand_labor-wage-dispute)	0.344554
(termlimit, tensys_strict)	0.330994

pairs	
(checks, xconst)	0.329545
(gq, gqi)	0.327936
(yrsoffc, maj)	0.316189
(totalseats, maj)	0.310318
(yrsoffc, stabs_strict)	0.302692
(military, xconst)	0.271173
(xconst, gq)	0.255725
(demand_price-increases/tax-policy, demand_political-behavior/process)	0.250856
(yrsoffc, checks)	0.244412
(termlimit, numvote)	0.235820
(numvote, tensys_strict)	0.234276
(demand_political-behavior/process, demand_police-brutality)	0.233095
(military, tensys_strict)	0.232767
(defmin, tensys_strict)	0.229613
(military, defmin)	0.224279
(military, checks)	0.224193
(checks, defmin)	0.222042
(maj, xconst)	0.221840
(termlimit, maj)	0.217684
(xconst, yrsoffc)	0.211793
(execnat, yrsoffc)	0.209019
(military, gq)	0.202514

## Modeling

Given the cleaned and aggregated dataset above, the next section moves into the Modeling phase. Each model type is constructed using elements of encoding, scaling, resampling and hyperparameter optimization.

- One hot encoding was essential given the categorical type of some features
- Standard scaling was essential given the vast array of different numerical feature distributions and ranges. Min-max scaling was considered but proved less effective.

- SMOTE was determined to be essential given the imbalanced nature of the dataset. Only 11% of the target feature values were 1, leaving the other 89% as 0. This is a prime example of the need for resampling, and SMOTE proved highly effective.
- Hyperparameter grid searches are inherently valuable when optimizing a model. Appropriate hyperparameter searches were used for each model type.

The output of each model is provided in terms of four core statistical measures (f1 score, accuracy, precision, and recall), in addition to displaying a confusion matrix for the test data. F1 was selected before the modeling process as the most relevant metric given that it encompasses all possible outcomes, as opposed to the other three metrics which leave out at least one possible outcome from their evaluation.

## Train-test split

```
In [21]: x_train, x_test, y_train, y_test = train_test_split(model_inputs, target,
                                                         random_state=RANDOM_STATE,
                                                         test_size=0.3)
```

## Define models and parameter grids

Define all models and grids in one place. A pipeline structure is created such that each of these models can be run with the below-defined hyperparameter tuning grids alongside their resampling, scaling and encoding. This allows for minimal repetition in code and a consistent structure.

```
In [22]: # Set parameter grid to search across
grid_bay = {'model__var_smoothing': [1e-9]}

grid_log = {'model__C': np.logspace(-1, 5, 10)}

grid_dt = {
    'model__max_depth': [3, 5, 7],
    'model__criterion': ['gini', 'entropy'],
    'model__min_samples_split': [5, 10],
    'model__min_samples_leaf': [5, 10]}

grid_rf = {
    'model__n_estimators': [25, 75, 150],
    'model__criterion': ['gini', 'entropy'],
    'model__max_depth': [3, 6, 10],
    'model__min_samples_split': [5, 10],
    'model__min_samples_leaf': [3, 6]}

grid_knn = {
    'model__leaf_size': [25, 50, 75],
    'model__n_neighbors': [3, 5, 7, 9],
    'model__weights': ['uniform', 'distance']}

grid_ada = {
    'model__n_estimators': [50, 200],
    'model__learning_rate': [1, 0.5, 0.25]}

grid_xgb = {
    'model__learning_rate': [0.1, 0.25, 0.5], #0.25
```

```
'model__max_depth': [6, 8, 10], #8
'model__n_estimators': [50, 100, 150, 200]} #100
```

```
np.random.seed(RANDOM_STATE)
model_bay = GaussianNB()
model_log = LogisticRegression(max_iter=5000)
model_dt = DecisionTreeClassifier()
model_rf = RandomForestClassifier()
model_knn = KNeighborsClassifier()
model_ada = AdaBoostClassifier(random_state=RANDOM_STATE)
model_xgb = XGBClassifier(eval_metric='logloss', use_label_encoder=False,
                          random_state=RANDOM_STATE)

grids = [grid_bay, grid_log, grid_dt, grid_rf, grid_knn, grid_ada, grid_xgb]
models = [model_bay, model_log, model_dt, model_rf, model_knn, model_ada, model_xgb]
```

## Pipeline function

This high-level function wraps all the different components of the model pipeline into one location, applying one-hot encoding, standard scaling, smote resampling, and grid searches to the input model. It also outputs performance in the form of standard metrics and a confusion matrix.

In [23]:

```
def create_pipeline_and_run(model, grid, metric='accuracy'):
    np.random.seed(RANDOM_STATE)
    ohe = OneHotEncoder(handle_unknown='ignore', sparse=False)
    scaler = StandardScaler()
    smote = SMOTE(random_state=RANDOM_STATE)

    selector_object = make_column_selector(dtype_exclude='number')
    selector_numeric = make_column_selector(dtype_include='number')
    transformer = make_column_transformer((ohe, selector_object),
                                          (scaler, selector_numeric))

    pipe = Pipeline([('transformer', transformer),
                     ('smote', smote),
                     ('model', model)])

    # Instantiate and fit grid search object
    grid = GridSearchCV(pipe, grid, scoring='f1', cv=10)
    grid.fit(x_train, y_train.values.ravel())
    pred = grid.best_estimator_.predict(x_test)

    print(f'{model}:')
    print_scores(pred, y_test)

    # Confusion matrix
    plt.figure()
    plot_confusion_matrix(grid.best_estimator_, x_test, y_test)
    plt.show();

    return grid.best_estimator_
```

## Dummy classifier as performance baseline



```
In [24]: for strategy in ["stratified", "uniform", "most_frequent"]:
        dummy_clf = DummyClassifier(strategy=strategy)
        dummy_clf.fit(x_train, y_train)

        print(f'DUMMY SCORE ({strategy}):')
        pred = dummy_clf.predict(x_test)
        print_scores(pred, y_test)
```

```
DUMMY SCORE (stratified):
- f1: 0.11267605633802816
- accuracy: 0.8226351351351351
- precision: 0.11299435028248588
- recall: 0.11235955056179775
DUMMY SCORE (uniform):
- f1: 0.17419060647514822
- accuracy: 0.4901463963963964
- precision: 0.10397387044093631
- recall: 0.5365168539325843
DUMMY SCORE (most_frequent):
- f1: 0.0
- accuracy: 0.8997747747747747
- precision: 0.0
- recall: 0.0
```


Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero\_division` parameter to control this behavior.

## Run *all models* defined above

Run this cell to output the performance of all above-defined models in one place for a side-by-side comparison

```
In [25]: pipes = []
        for grid, model in zip(grids, models):
            pipe = create_pipeline_and_run(model, grid)
            pipes.append(pipe)
```

```
GaussianNB():
- f1: 0.24158696422245837
- accuracy: 0.39724099099099097
- precision: 0.13822456424807458
- recall: 0.9578651685393258
<Figure size 432x288 with 0 Axes>
```




```
LogisticRegression(max_iter=5000):
```

```

- f1: 0.3941818181818182
- accuracy: 0.7654842342342343
- precision: 0.26594700686947986
- recall: 0.7612359550561798
<Figure size 432x288 with 0 Axes>


```



```

DecisionTreeClassifier():
- f1: 0.4162361623616236
- accuracy: 0.7773085585585585
- precision: 0.2822822822822823
- recall: 0.7921348314606742
<Figure size 432x288 with 0 Axes>


```




```

RandomForestClassifier():
- f1: 0.47534516765286
- accuracy: 0.8502252252252253
- precision: 0.3662613981762918
- recall: 0.6769662921348315
<Figure size 432x288 with 0 Axes>


```



```
KNeighborsClassifier():
- f1: 0.5446716899892357
- accuracy: 0.8809121621621622
- precision: 0.44153577661431065
- recall: 0.7106741573033708
<Figure size 432x288 with 0 Axes>
```




```
AdaBoostClassifier(random_state=2021):
- f1: 0.3880597014925373
- accuracy: 0.8499436936936937
- precision: 0.32815533980582523
- recall: 0.4747191011235955
<Figure size 432x288 with 0 Axes>
```



```
XGBClassifier(base_score=None, booster=None, colsample_bylevel=None,
              colsample_bynode=None, colsample_bytree=None,
              eval_metric='logloss', gamma=None, gpu_id=None,
              importance_type='gain', interaction_constraints=None,
              learning_rate=None, max_delta_step=None, max_depth=None,
              min_child_weight=None, missing=nan, monotone_constraints=None,
              n_estimators=100, n_jobs=None, num_parallel_tree=None,
              random_state=2021, reg_alpha=None, reg_lambda=None,
              scale_pos_weight=None, subsample=None, tree_method=None,
              use_label_encoder=False, validate_parameters=None,
              verbosity=None):
```

```
- f1: 0.7769347496206372
- accuracy: 0.9586148648648649
- precision: 0.8448844884488449
- recall: 0.7191011235955056
<Figure size 432x288 with 0 Axes>
```



## Run *only one* model

Choose which model to run in the below cell (this cell is only used for iterative testing and investigating model specifics without running all models)

```
In [26]: # Uncomment and run to look at one model separately
#xgb = create_pipeline_and_run(model_xgb, grid_xgb);
```

### Print optimal model hyperparameters

```
In [27]: xgb = pipes[-1] # Since it is the last model in pipes
print(xgb.steps[2])

('model', XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                        colsample_bynode=1, colsample_bytree=1, eval_metric='logloss',
                        gamma=0, gpu_id=-1, importance_type='gain',
                        interaction_constraints='', learning_rate=0.5, max_delta_step=0,
                        max_depth=10, min_child_weight=1, missing=nan,
                        monotone_constraints='()', n_estimators=200, n_jobs=8,
                        num_parallel_tree=1, random_state=2021, reg_alpha=0, reg_lambda=1,
                        scale_pos_weight=1, subsample=1, tree_method='exact',
                        use_label_encoder=False, validate_parameters=1, verbosity=None))
```

## Test model on test dataset


XG boost proves to be the highest performing model. Test its performance on the test dataset.

In [28]:

```
# Predict output
pred = xgb.predict(x_test)

# Show performance
print_scores(pred, y_test)
plot_confusion_matrix(xgb, x_test, y_test);
```

```
- f1: 0.7769347496206372
- accuracy: 0.9586148648648649
- precision: 0.8448844884488449
- recall: 0.7191011235955056
```



### Plot confusion matrices

In [29]:

```
# Predict output
pred = xgb.predict(x_test)

# Show performance
print_scores(pred, y_test)

# Plot test data and full data performance
labels = ['No change', 'Regime Change']
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(12,8))
plt.subplots_adjust(wspace=0.6, hspace=None)
axes[0].set_title('Test Dataset (%)')
axes[1].set_title('Entire Dataset (%)')

plot_confusion_matrix(xgb, x_test, y_test,
                      ax=axes[0],
                      display_labels=labels,
                      colorbar=False,
                      normalize='all')

plot_confusion_matrix(xgb, model_inputs, target,
                      ax=axes[1],
                      display_labels=labels,
                      colorbar=False,
                      normalize='all');

plt.savefig('../images/confusion_matrices.png')
```

```
- f1: 0.7769347496206372
```

- accuracy: 0.9586148648648649
- precision: 0.8448844884488449
- recall: 0.7191011235955056




## Feature importance

Evaluate the feature importance in the top-performing model.

In [30]:

```
# SHAP summary plot for XGB
produce_shap_plot(x_train, y_train, x_test, y_test, clone(xgb),
                 title='Feature Importance Summary Plot for XGB Model',
                 savepath = '../images/shap_summary_plot.png');
```



In [31]:

```
# SHAP bar plot for XGB model
x_tr_manual, y_tr_manual, x_te_manual = get_shap_df(x_train, y_train, x_test)
```


```

model = xgb.steps[2][1]

# Calculate SHAP values
explainer = shap.Explainer(model)
shap_values = explainer(x_te_manual)

# SHAP bar plot for XGB model
plt.figure()
plt.title('Feature Importance Bar Plot for XGB Model', fontsize=20)
shap.plots.bar(shap_values, max_display=20, show=False)
plt.savefig('../images/shap_bar_plot.png');

```



## Permutation Feature Importance

Source: [https://scikit-learn.org/stable/modules/permutation\\_importance.html](https://scikit-learn.org/stable/modules/permutation_importance.html)

```

In [32]: # Subsequent functions can't handle categoricals so
          # fit dataframe outside pipeline

```

```

model = xgb.steps[2][1]
x_tr, y_tr, x_te = get_shap_df(x_train, y_train, x_test)
model.fit(x_tr, y_tr);

r = permutation_importance(model, x_te, y_test,
                           n_repeats=30,
                           random_state=RANDOM_STATE)

for i in r.importances_mean.argsort()[::-1]:
    if r.importances_mean[i] - 3 * r.importances_std[i] > 0:
        #print(f"{df_train.feature_names[i]:<8}")
        print(f"{x_tr.columns[i]} "
              f"{r.importances_mean[i]:.3f}")

```

```

xconst 0.065
tensys_strict 0.023
yrsoffc 0.020
totalseats 0.018
protestnumber_log 0.016
maj 0.013
herfgov 0.007
numvote 0.006
gq 0.004
region_Europe 0.003
military 0.002
country_Ethiopia 0.001
region_MENA 0.001
country_Niger 0.001
country_Guyana 0.001
execrel_Islamic 0.001
execrel_OTHER 0.001
country_Haiti 0.000
country_Guinea 0.000
country_Venezuela 0.000
country_Paraguay 0.000
country_Comoros 0.000
country_Turkey 0.000
country_Sierra Leone 0.000
country_Senegal 0.000

```

## Export final datasets and models for use in other notebooks

The final, cleaned dataframe with features selected and engineered is exported via SQL to be used in the EDA notebook. Three selected final, fitted models are stored as Pickle files for use in the Final Presentation notebook.

### Export data via SQL

In [33]:

```

engine = create_engine('sqlite:///../data/processed/all_data.db')

model_data = pd.concat([model_inputs, target], axis=1)

with engine.begin() as connection:
    model_data.to_sql(name='all_modeled_data',
                      con=connection,

```



```
if_exists='replace',  
index=False)
```

## Export models via Pickle

In [34]:

```
# Save fitted Logistic Regression  
with open('../data/processed/model_logreg.pickle', 'wb') as f:  
    pickle.dump(pipes[1], f)  
  
# Save fitted Random Forest  
with open('../data/processed/model_rf.pickle', 'wb') as f:  
    pickle.dump(pipes[3], f)  
  
# # Save fitted XG Boost  
with open('../data/processed/model_xgb.pickle', 'wb') as f:  
    pickle.dump(pipes[-1], f)
```